
pybedtools Documentation

Release 0.2.3dev

Ryan Dale

May 22, 2011

CONTENTS

1	Overview	1
2	Contents:	3
2.1	Installation	3
2.2	Three brief examples	4
2.3	Tutorial Contents	8
2.4	Topical Documentation	19
2.5	pybedtools Reference	32
3	Indices and tables	57
	Python Module Index	59
	Index	61

OVERVIEW

pybedtools is a Python wrapper for Aaron Quinlan's BEDtools and is designed to leverage the “genome algebra” power of BEDtools from within Python scripts.

This documentation is written assuming you know how to use BEDTools and Python.

See full online documentation, including installation instructions, at <http://pybedtools.genomicnorth.com>.

Note: `pybedtools` is still very much in progress. Please keep that in mind when assessing whether to use this package in production code.

The documentation is separated into 4 main parts, depending on the depth you'd like to cover. Lazy, or just want to jump in? Check out *Three brief examples* to get a feel for the package. Want a guided tour? Give the *Tutorial Contents* a shot. More advanced features are described in the *Topical Documentation* section. Finally, doctested module documentation can be found in *pybedtools Reference*.

CONTENTS:

2.1 Installation

2.1.1 Python requirements

- Python 2.5 or greater; Python 3 support coming soon
- `argparse` if you are running Python < 2.7 (Python 2.7 comes with `argparse` already)
- `Cython` - part of `pybedtools` is written in `Cython` for speed

Both `argparse` and `Cython` can be installed with `pip`:

```
pip install cython argparse
```

or `easy_install`:

```
easy_install cython argparse
```

2.1.2 Installing `pybedtools`

To install the latest version of `pybedtools` you have 2 options:

Option 1: install from source

- go to <http://github.com/daler/pybedtools>
- click the Downloads link (**ldll**)
- choose either a `.tar.gz` or a `.zip` file, whatever you're comfortable with
- unzip into a temporary directory
- from the command line, run:

```
python setup.py install
```

(you may need admin rights to do this)

Option 2: use `pip` to automatically download the latest stable version from the [Python Package Index](#):

```
sudo pip install --upgrade pybedtools
```

Warning: These docs are written for the latest version on git. For docs specific to the version you have installed, please see the docs included with that version

2.1.3 Installing BEDTools_

`pybedtools` relies heavily on `BEDTools_`. To install `BEDTools_`,

- follow the instructions at <https://github.com/ark5x/bedtools> to install
- make sure all its programs are on your path

2.1.4 Testing your installation

A quick functional test is to create a new script with the following contents:

```
import pybedtools
a = pybedtools.example_bedtool('a.bed')
b = pybedtools.example_bedtool('b.bed')
print a.intersect(b)
```

If this script is called `test.py`, then running it with `python test.py` should print out:

```
chr1      155      200      feature2      0      +
chr1      155      200      feature3      0      -
chr1      900      901      feature4      0      +
```

For more extensive testing:

- If you have `nosetest` installed (e.g., via `pip install nose`) you can run the test suite with:

```
sh test.sh $VERSION
```

where `$VERSION` is the version of Python you'd like to run the tests with, e.g., `2.7`.

- If you have `sphinx` installed (e.g., via `pip install sphinx`), you can run the doctests by going to the `docs` directory and running:

```
make doctest
```

2.2 Three brief examples

Here are three examples to show typical usage of `pybedtools`. More info can be found in the docstrings of `pybedtools` methods and in the *Tutorial Contents*.

2.2.1 Example 1: Save a BED file of intersections, with track line

This example saves a new BED file of intersections between `a.bed` and `b.bed`, adding a track line to the output:

```
>>> import pybedtools
>>> a = pybedtools.BedTool('a.bed')
>>> a.intersect('b.bed').saveas('a-and-b.bed', trackline="track name='a and b' color=128,0,0")
```

2.2.2 Example 2: Intersections for a 3-way Venn diagram

This example gets values for a 3-way Venn diagram of overlaps. This demonstrates operator overloading of `BedTool` objects:


```
>>> import pybedtools

>>> # set up 3 different bedtools
>>> a = pybedtools.BedTool('a.bed')
>>> b = pybedtools.BedTool('b.bed')
>>> c = pybedtools.BedTool('c.bed')

>>> (a-b-c).count() # unique to a
>>> (a+b-c).count() # in a and b, not c
>>> (a+b+c).count() # common to all
>>> # ... and so on, for all the combinations.
```

2.2.3 Example 3: Classify reads into intron/exon

This example is a little more involved, but highlights some useful features of `pybedtools`. Here, we classify reads in a BAM file into exon, intron, exon and intron (i.e., exonic in one isoform but intronic in another overlapping isoform) or intergenic classes in a stranded manner. A more generalized version of this example can be found in the `pybedtools.scripts.classify_reads` script.

```
import pybedtools

# Create a BedTool for the GFF file of annotations
g = pybedtools.BedTool('example.gff')

# Set up two functions that will filter and then rename features to set up for
# merging

def renamer(x):
    """
    *x* is an Interval object representing a GFF feature.

    Renames the feature after the feature type; this is needed for when
    .merge() combines names together in a later step.
    """

    # This illustrates setting and getting fields in an Interval object based
    # on attribute or index
    x.name = x[2]
    return x

def filter_func(x):
    """
    *x* is an Interval object representing a GFF feature.

    This filter function will only pass features of type "intron" or "exon"
    """
    if x[2] in ('intron', 'exon'):
        return True
    return False

# Filter and rename the GFF features by passing the above functions to
# .filter() and .each(). Note that since each method returns a new BedTool,
# methods can be chained together
```

```
g2 = g.filter(filter_func).each(renamer)

# Save a copy of the new GFF file for later inspection
g2 = g2.saveas('edited.gff')

# Here we call mergeBed, which operates on the file pointed to by g2
# (that is, 'edited.gff').
#
# We use several options for BEDTools mergeBed:
#
# 'nms' combines names of merged features (after filtering and renaming, this
# is either "intron" or "exon") into a semicolon-delimited list;
#
# d=-1 does not merge bookended features together;
#
# s=True ensures a stranded merge;
#
# scores='sum' ensures a valid BED file result, with a score field before the
# strand field
#
merged = g2.merge(nms=True, d=-1, s=False, scores='sum')

# Next, we intersect a BAM file with the merged features. Here, we explicitly
# specify the 'abam' and 'b' arguments, ensure stranded intersections, use
# BED-format output, and report the entire a and b features in the output:
#
reads_in_features = merged.intersect(abam='example.bam',
                                     b=merged.fn,
                                     s=True,
                                     bed=True,
                                     wao=True)

# Set up a dictionary to hold counts
from collections import defaultdict
results = defaultdict(int)

# Iterate through the intersected reads, parse out the names of the features
# they intersected, and increment counts in the dictionary. This illustrates
# how BedTool objects follow the iterator protocol, each time yielding an
# Interval object:
#
total = 0.0
for intersected_read in reads_in_features:
    total += 1

    # Extract the name of the feature this read intersected by indexing into
    # the Interval
    intersected_feature = feature[-4]

    # Convert names like "intron;intron;intron", which indicates overlapping
    # isoforms or genes all with introns in this region, to the simple class of
    # "intron"
    key = ';'.join(sorted(list(set(intersected_with.split(';')))))

    # Increment the count for this class
    results[key] += 1
```

```
# Rename the "." key to something more meaningful
results['intergenic'] = results.pop('.')

# Add the total to the dictionary
results['total'] = int(total)

print results

# Delete any temporary files created
pybedtools.cleanup()
```

Here is the same script without comments, to give a better feel for `pybedtools`:

```
import pybedtools

g = pybedtools.BedTool('example.gff')

def renamer(x):
    x.name = x[2]
    return x

def filter_func(x):
    if x[2] in ('intron', 'exon'):
        return True
    return False

g2 = g.filter(filter_func).each(renamer)
g2 = g2.saveas('edited.gff')
merged = g2.merge(nms=True, d=-1, s=False, scores='sum')
reads_in_features = merged.intersect(abam='example.bam',
                                     b=merged.fn,
                                     s=True,
                                     bed=True,
                                     wao=True)

from collections import defaultdict
results = defaultdict(int)
total = 0.0
for intersected_read in reads_in_features:
    total += 1
    intersected_feature = feature[-4]
    key = ';'.join(sorted(list(set(intersected_with.split(';')))))
    results[key] += 1

results['intergenic'] = results.pop('.')
results['total'] = int(total)
print results

pybedtools.cleanup()
```

For more, continue on to the [Tutorial Contents](#), and then check out the [Topical Documentation](#).

2.3 Tutorial Contents

2.3.1 Intro

This tutorial assumes that

1. You know how to use [BEDTools](#) (if not, check out the [BEDTools documentation](#))
2. You know how to use Python (if not, check out some tutorials like [Learn Python the Hard Way](#))

A brief note on conventions

Throughout this documentation I've tried to use consistent typography, as follows:

- Python variables and arguments, as well as filenames look like this: `s=True`
- Methods, which are often linked to documentation look like this: `BedTool.merge()`.
- Arguments that are passed to [BEDTools](#) programs, as if you were on the command line, look like this: `-d`.
- The “>>>” in the examples below indicates a Python interpreter prompt and means to type the code into an interactive Python interpreter like [IPython](#) or in a script. (don't type the >>>)

Onward!

2.3.2 Create a BedTool

First, follow the [Installation](#) instructions if you haven't already done so to install both [BEDTools](#) and `pybedtools`.

Then import the `pybedtools` module:

```
>>> import pybedtools
```

This documentation will use example files that ship with `pybedtools`, but you may find it more useful to use your own files. To create a `BedTool` you need to specify a filename. To get a list of example files that ship with `pybedtools`:

```
>>> pybedtools.list_example_files()
['a.bed', 'b.bed', 'c.gff', 'd.gff', 'dm3-chr2L-5M-invalid.gff.gz', 'dm3-chr2L-5M.gff.gz', 'dmel-all
```

Note that there are BED and GFF files, some of which are compressed. All files supported by [BEDTools](#) are supported by `pybedtools`.

Using one of these filenames, get the full path to the file (which will depend on your operating system and how you installed the package) with:

```
>>> full_path = pybedtools.example_filename('a.bed')
```

Once you have a filename (whether it's an example file or your own file), creating a `BedTool` is easy:

```
>>> # create a new BedTool using that filename
>>> a = pybedtools.BedTool(full_path)
```

Let's set up a second `BedTool` so we can do intersections and subtractions. This time, we'll use a convenience function for creating `BedTool` instances out of example files (if you're using your own files, just make another one the same way `a` was made above).

```
>>> # create another BedTool to play around with
>>> b = pybedtools.example_bedtool('b.bed')
```

See [Creating a BedTool](#) for more information, including making `BedTool` objects directly from strings and iterators.

2.3.3 Intersections

One common use of **BEDTools** and `pybedtools` is to perform intersections.

First, create some example `BedTool` instances if you haven't already done so:

```
>>> a = pybedtools.example_bedtool('a.bed')
>>> b = pybedtools.example_bedtool('b.bed')
```

Then do the intersection with the `BedTool.intersect()` method:

```
>>> a_and_b = a.intersect(b)
```

`a_and_b` is a new `BedTool` instance. It now points to a temp file on disk, which is stored in the attribute `a_and_b.fn`; this temp file contains the intersection of `a` and `b`.

We can either print the new `BedTool` (which will show ALL features – use with caution if you have huge files!) or use the `BedTool.head()` method to show up to the first N lines (10 by default). Here's what `a`, `b`, and `a_and_b` look like:

```
>>> a.head()
chr1    1    100 feature1    0    +
chr1    100 200 feature2    0    +
chr1    150 500 feature3    0    -
chr1    900 950 feature4    0    +

>>> b.head()
chr1    155 200 feature5    0    -
chr1    800 901 feature6    0    +

>>> a_and_b.head()
chr1    155 200 feature2    0    +
chr1    155 200 feature3    0    -
chr1    900 901 feature4    0    +
```

The `BedTool.intersect()` method simply wraps the **BEDTools** program `intersectBed`. This means that we can pass `BedTool.intersect()` any arguments that `intersectBed` accepts. For example, if we want to use the `intersectBed` switch `-u` (which acts as a True/False switch to indicate that we want to see the features in `a` that overlapped something in `b`), then we can use the keyword argument `u=True`, like this:

```
>>> # Intersection using the -u switch
>>> a_with_b = a.intersect(b, u=True)
>>> a_with_b.head()
chr1    100 200 feature2    0    +
chr1    150 500 feature3    0    -
chr1    900 950 feature4    0    +
```

This time, `a_with_b` is another `BedTool` object that points to a different temp file whose name is stored in `a_with_b.fn`. You can read more about the use of temp files in [Principle 1: Temporary files are created automatically](#). More on arguments that you can pass to `BedTool` objects in a moment, but first, some info about saving files.

2.3.4 Saving the results

If you want to save the results as a meaningful filename for later use, use the `BedTool.saveas()` method. This also lets you optionally specify a trackline for directly uploading to the UCSC Genome Browser, instead of opening

up the files afterward and manually adding a trackline:

```
>>> c = a_with_b.saveas('intersection-of-a-and-b.bed', trackline='track name="a and b"')
>>> print c.fn
intersection-of-a-and-b.bed

>>> print c
track name="a and b"
chr1      155      200      feature2      0      +
chr1      155      200      feature3      0      -
chr1      900      901      feature4      0      +
```

Note that the `BedTool.saveas()` method returns a new `BedTool` object which points to the newly created file on disk. This allows you to insert a `BedTool.saveas()` call in the middle of a chain of commands (described in another section below).

2.3.5 Default arguments

Recall that we passed the `u=True` argument to `a.intersect()`:

```
>>> a_with_b = a.intersect(b, u=True)
```

While we're on the subject of arguments, note that we didn't have to specify `-a` or `-b` arguments, like you would need if calling `intersectBed` from the command line. That's because `BedTool` objects make some assumptions for convenience.

We could have supplied the arguments `a=a.fn` and `b=b.fn`:

```
>>> another_way = a.intersect(a=a.fn, b=b.fn, u=True)
>>> another_way == a_with_b
True
```

But since we're calling a method on the `BedTool` object `a`, `pybedtools` assumes that the file `a` points to (stored in the attribute `a.fn`) is the one we want to use as input. So by default, we don't need to explicitly give the keyword argument `a=a.fn` because the `a.intersect()` method does so automatically.

We're also calling a method that takes a second bed file as input – other such methods include `BedTool.subtract()` and `BedTool.closest()`, and others. For these methods, in addition to assuming `-a` is taken care of by the `BedTool.fn` attribute, `pybedtools` also assumes the first unnamed argument to these methods are the second file you want to operate on (and if you pass a `BedTool`, it'll automatically use the file in the `fn` attribute of that `BedTool`).

An example may help to illustrate: these different ways of calling `BedTool.intersect()` all have the same results, with the first version being the most compact (and probably most convenient):

```
>>> # these all have identical results
>>> x1 = a.intersect(b)
>>> x2 = a.intersect(a=a.fn, b=b.fn)
>>> x3 = a.intersect(b=b.fn)
>>> x4 = a.intersect(b, a=a.fn)
>>> x1 == x2 == x3 == x4
True
```

Note that `a.intersect(a=a.fn, b)` is not a valid Python expression, since non-keyword arguments must come before keyword arguments, but `a.intersect(b, a=a.fn)` works fine.

If you're ever unsure, the docstring for these methods indicates which, if any, arguments are used as default. For example, in the `BedTool.intersect()` help, it says:

.. note::

For convenience, the file this bedtool object points to is passed as "-a"

OK, enough about arguments for now, but you can read more about them in *Principle 2: Names and arguments are as similar as possible to BEDTools*, *Principle 4: Sensible default args* and *Principle 5: Other arguments have no defaults*.

2.3.6 Chaining methods together (pipe)

One useful thing about BedTool methods is that they often return a new BedTool. In practice, this means that we can chain together multiple method calls all in one line, similar to piping on the command line.

For example, this intersect and merge can be combined into one command:

```
>>> # These two lines...
>>> x1 = a.intersect(b, u=True)
>>> x2 = x1.merge()

>>> # ...can be combined into one line:
>>> x3 = a.intersect(b, u=True).merge()

>>> x2 == x3
True
```

In general, methods that return BedTool objects have the following text in their docstring to indicate this:

.. note::

This method returns a new BedTool instance

A rule of thumb is that all methods that wrap **BEDTools** programs return BedTool objects, so you can chain these together. Many pybedtools-unique methods return BedTool objects too, just check the docs (according to *Principle 7: Check the help*). For example, as we saw in one of the examples above, the `BedTool.saveas()` method returns a BedTool object. That means we can sprinkle those commands within the example above to save the intermediate steps as meaningful filenames for later use. For example:

```
>>> x4 = a.intersect(b, u=True).saveas('a-with-b.bed').merge().saveas('a-with-b-merged.bed')
```

Now we have new files in the current directory called `a-with-b.bed` and `a-with-b-merged.bed`. Since `BedTool.saveas()` returns a BedTool object, `x4` points to the `a-with-b-merged.bed` file.

2.3.7 Operator overloading

There's an even easier way to chain together commands.

I found myself doing intersections so much that I thought it would be useful to overload the + and - operators to do intersections. To illustrate, these two example commands do the same thing:

```
>>> x5 = a.intersect(b, u=True)
>>> x6 = a + b

>>> x5 == x6
True
```

Just as the `+` operator assumes `intersectBed` with the `-u` arg, the `-` operator assumes `intersectBed` with the `-v` arg:

```
>>> x7 = a.intersect(b, v=True)
>>> x8 = a - b

>>> x7 == x8
True
```

If you want to operating on the resulting `BedTool` that is returned by an addition or subtraction, you'll need to wrap the operation in parentheses. This is another way to do the chaining together of the intersection and merge example from above:

```
>>> x9 = (a + b).merge()
```

And to double-check that all these methods return the same thing:

```
>>> x2 == x3 == x4 == x9
True
```

You can learn more about chaining in *Principle 6: Chaining together commands*.

2.3.8 Intervals

An `Interval` object is how `pybedtools` represents a line in a BED, GFF, GTF, or VCF file in a uniform fashion. This section will describe some useful features of `Interval` objects.

First, let's get a `BedTool` to work with:

```
>>> a = pybedtools.example_bedtool('a.bed')
```

We can access the `Intervals` of a several different ways. The most common use is probably by using the `BedTool` `a` as an iterator. For now though, let's look at a single `Interval`:

```
>>> feature = iter(a).next()
```

Common `Interval` attributes

Printing a feature converts it into the original line from the file:

```
>>> print feature
chr1      1      100      feature1      0      +
```

All features have `chrom`, `start`, `stop`, `name`, `score`, and `strand` attributes. Note that `start` and `stop` are long integers, while everything else (including `score`) is a string.

```
>>> feature.chrom
'chr1'
```

```
>>> feature.start
1L
```

```
>>> feature.stop
100L
```

```
>>> feature.name
'feature1'
```



```
>>> feature.score
'0'
```

```
>>> feature.strand
'+'
```

Let's make another feature that only has chrom, start, and stop to see how `pybedtools` deals with missing attributes:

```
>>> feature2 = iter(pybedtools.BedTool('chrX 500 1000', from_string=True)).next()
```

```
>>> print feature2
chrX      500      1000
```

```
>>> feature2.chrom
'chrX'
```

```
>>> feature2.start
500L
```

```
>>> feature2.stop
1000L
```

```
>>> feature2.name
''
```

```
>>> feature2.score
''
```

```
>>> feature2.strand
''
```

This illustrates that default values are empty strings.

Indexing into `Interval` objects

`Interval` objects can also be indexed by position into the original line (like a list) or indexed by name of attribute (like a dictionary).

```
>>> print feature
chr1      1      100      feature1      0      +
```

```
>>> feature[0]
'chr1'
```

```
>>> feature['chrom']
'chr1'
```

```
>>> feature[1]
'1'
```

```
>>> feature['start']
1L
```

Fields

Interval objects have a `Interval.fields` attribute that contains the original line split into a list of strings. When an integer index is used on the Interval (for example, `feature[3]`), it is the `fields` attribute that is actually being indexed into.

```
>>> f = iter(pybedtools.BedTool('chr1 1 100 asdf 0 + a b c d', from_string=True)).next()
>>> f.fields
['chr1', '1', '100', 'asdf', '0', '+', 'a', 'b', 'c', 'd']
>>> len(f.fields)
10
```

BED is 0-based, others are 1-based

One troublesome part about working with multiple formats is that BED files have a different coordinate system than GFF/GTF/VCF/ files.

BED files are 0-based (the first base of the chromosome is considered position 0) and the **feature does not include the stop position**.

GFF, GTF, and VCF files are 1-based (the first base of the chromosome is considered position 1) and the **feature includes the stop position**.

Note: `pybedtools` follows the following conventions:

- The value in `Interval.start` will **always** contain the 0-based start position, even if it came from a GFF or other 1-based feature.
 - Getting the `len()` of an Interval will always return `Interval.stop - Interval.start`, so no matter what format the original file was in, the length will be correct.
 - The contents of `Interval.fields` will **always** be strings, which in turn always represent the original line in the file. This means that for a GFF feature, `Interval.fields[3]` or `Interval[3]`, which is 1-based according to the file format, will always be one bp larger than `Interval.start`, which always contains the 0-based start position. However, `Interval[3]` will be a string and `Interval.start` will be a long.
-

To illustrate and confirm, let's create a GFF feature and a BED feature from scratch and compare them:

```
>>> # GFF Interval from scratch
>>> gff = ["chr1",
...       "fake",
...       "mRNA",
...       "51",    # <- start is 1 greater than start for the BED feature below
...       "300",
...       ".",
...       "+",
...       ".",
...       "ID=mRNA1;Parent=genel;"]
>>> gff = pybedtools.create_interval_from_list(gff)
>>> print gff
chr1      fake      mRNA      51      300      .      +      .      ID=mRNA1;Parent=genel;

>>> # BED Interval from scratch
>>> bed = ["chr1",
...       "50",
...       "300",
...       "mRNA1",
```

```

...         ".",
...         "+"]
>>> bed = pybedtools.create_interval_from_list (bed)
>>> print bed
chr1      50      300      mRNA1      .      +

>>> # confirm they are recognized as the right type
>>> gff.file_type
'gff'
>>> bed.file_type
'bed'

>>> # Start attributes should be identical
>>> bed.start == gff.start == 50
True

>>> bed.start
50L
>>> bed[1]
'50'

>>> # GFF .start is 1 less than the string value stored at index 3
>>> gff.start
50L
>>> gff[3]
'51'

>>> len(bed) == len(gff) == 250
True

```

GFF features have access to attributes

GFF and GTF files have lots of useful information in their attributes field (the last field in each line). These attributes can be accessed with the `Interval.attrs` attribute, which acts like a dictionary. For speed, the attributes are lazy – they are only parsed when you ask for them. BED files, which do not have an attributes field, will return an empty dictionary.

```

>>> # original feature
>>> print gff
chr1      fake      mRNA      51      300      .      +      .      ID=mRNA1;Parent=genel;

>>> # original attributes
>>> gff.attrs
{'ID': 'mRNA1', 'Parent': 'genel'}

>>> # add some new attributes
>>> gff.attrs['Awesomeness'] = 99
>>> gff.attrs['ID'] = 'transcript1'

>>> # Changes in attributes are propagated to the printable feature
>>> print gff
chr1      fake      mRNA      51      300      .      +      .      Awesomeness=99;ID=transcript1;Parent=genel;

```

Understanding `Interval` objects is important for using the powerful filtering and mapping facilities of `BedTool` objects, as described in the next section.

2.3.9 Filtering

The `filter()` method lets you pass in a function that accepts an `Interval` as its first argument and returns `True` for `False`. For example, here's how to get a new `BedTool` containing features from `a` that are more than 100 bp long:

```
>>> a = pybedtools.example_bedtool('a.bed')
>>> b = a.filter(lambda x: len(x) > 100)
>>> print b
chr1      150      500      feature3      0      -
```

The `filter()` method will pass its `*args` and `**kwargs` to the function provided. So here is a more generic case would be the following, where the function is defined once and different arguments are passed in for filtering on different lengths:

```
>>> def len_filter(feature, L):
...     "Returns True if feature is longer than L"
...     return len(feature) > L
```

Now we can pass different lengths without defining a new function for each length of interest:

```
>>> a = pybedtools.example_bedtool('a.bed')

>>> print a.filter(len_filter, L=10)
chr1      1      100      feature1      0      +
chr1      100     200      feature2      0      +
chr1      150     500      feature3      0      -
chr1      900     950      feature4      0      +

>>> print a.filter(len_filter, L=99)
chr1      100     200      feature2      0      +
chr1      150     500      feature3      0      -

>>> print a.filter(len_filter, L=200)
chr1      150     500      feature3      0      -
```

See *Using BedTool objects as iterators/generators* for more advanced and space-efficient usage of `filter()` using iterators.

Fast filtering functions in Cython

The `featurefuncs` module contains some ready-made functions written in Cython that will be faster than pure Python equivalents. For example, there are `greater_than()` and `less_than()` functions, which are about 70% faster. In IPython:

```
>>> from pybedtools.featurefuncs import greater_than

>>> len(a)
310456

>>> def L(x,width=100):
...     return len(x) > 100

>>> %timeit a.filter(greater_than, 100)
1 loops, best of 3: 1.74 s per loop
```

```
>>> %timeit a.filter(L, 100)
1 loops, best of 3: 2.96 s per loop
```

2.3.10 Each

Similar to `BedTool.filter()`, which applies a function to return True or False given an Interval, the `BedTool.each()` method applies a function to return a new, possibly modified Interval.

The `BedTool.each()` method applies a function to every feature. Like `BedTool.filter()`, you can use your own function or some pre-defined ones in the `featurefuncs` module. Also like `filter()`, `*args` and `**kwargs` are sent to the function.

```
>>> a = pybedtools.example_bedtool('a.bed')
>>> b = pybedtools.example_bedtool('b.bed')

>>> # The results of an "intersect" with c=True will return features
>>> # with an additional field representing the counts.
>>> with_counts = a.intersect(b, c=True)
```

Let's define a function that will take the number of counts in each feature as calculated above and divide by the number of bases in that feature. We can also supply an optional scalar, like 0.001, to get the results in "number of intersections per kb". We then insert that value into the score field of the feature. Here's the function:

```
>>> def normalize_count(feature, scalar=0.001):
...     """
...     assume feature's last field is the count
...     """
...     counts = float(feature[-1])
...     normalized = counts / len(feature) * scalar
...
...     # need to convert back to string to insert into feature
...     feature.score = str(normalized)
...     return feature
```

And we apply it like this:

```
>>> normalized = with_counts.each(normalize_count)
>>> print normalized
chr1      1      100      feature1      0.0      +      0
chr1     100     200      feature2     1e-05      +      1
chr1     150     500      feature3     2.85714285714e-06      -      1
chr1     900     950      feature4     2e-05      +      1
```

2.3.11 Using the history and tags

BEDTools makes it very easy to do rather complex genomic algebra. Sometimes when you're doing some exploratory work, you'd like to rewind back to a previous step, or clean up temporary files that have been left on disk over the course of some experimentation.

To assist this sort of workflow, `BedTool` instances keep track of their history in the `BedTool.history` attribute. Let's make an example `BedTool`, `c`, that has some history:

```
>>> a = pybedtools.example_bedtool('a.bed')
>>> b = pybedtools.example_bedtool('b.bed')
>>> c = a.intersect(b, u=True)
```

c now has a history which tells you all sorts of useful things (described in more detail below):

```
>>> print c.history
[<HistoryStep> bedtool("/home/ryan/pybedtools/pybedtools/test/a.bed").intersect("/home/ryan/pybedtools/test/b.bed")]
```

There are several things to note here. First, the history describes the full commands, including all the names of the temp files and all the arguments that you would need to run in order to re-create it. Since `BedTool` objects are fundamentally file-based, the command refers to the underlying filenames (i.e., `a.bed` and `b.bed`) instead of the `BedTool` instances (i.e., `a` and `b`). A simple copy-paste of the command will be enough re-run the command. While this may be useful in some situations, be aware that if you do run the command again you'll get *another* temp file that has the same contents as `c`'s temp file.

To avoid such cluttering of your temp dir, the history also reports **tags**. `BedTool` objects, when created, get a random tag assigned to them. You can get the `BedTool` associated with tag with the `pybedtools.find_tagged()` function. These tags are used to keep track of instances during this session.

So in this case, we could get a reference to the `a` instance with:

```
>>> should_be_a = pybedtools.find_tagged('klkreuay')
```

Here's confirmation that the parent of the first step of `c`'s history is `a` (note that `HistoryStep` objects have a `HistoryStep.parent_tag` and `HistoryStep.result_tag`):

```
>>> pybedtools.find_tagged(c.history[0].parent_tag) == a
True
```

Let's make something with a more complicated history:

```
>>> a = pybedtools.example_bedtool('a.bed')
>>> b = pybedtools.example_bedtool('b.bed')
>>> c = a.intersect(b)
>>> d = c.slop(g=pybedtools.chromsizes('hg19'), b=1)
>>> e = d.merge()

>>> # this step adds complexity!
>>> f = e.subtract(b)
```

Let's see what the history of `f` (the last `BedTool` created) looks like . . . note that here I'm formatting the results to make it easier to see:

```
>>> print f.history
[
|   [
|   |   [
|   |   |   [
|   |   |   |<HistoryStep> BedTool("/usr/local/lib/python2.6/dist-packages/pybedtools/test/data/a.bed")
|   |   |   |                               "/usr/local/lib/python2.6/dist-packages/pybedtools/test/data/b.bed")
|   |   |   |                               ),
|   |   |   |                               parent tag: rzzrtxlv,
|   |   |   |                               result tag: ifbsanqk
|   |   |   ],
|   |   ],
|   |   |<HistoryStep> BedTool("/tmp/pybedtools.BgULVj.tmp").slop(
|   |   |                               b=1, genome="hg19"
|   |   |                               ),
|   |   |                               parent tag: ifbsanqk,
|   |   |                               result tag: omfrkwjp
|   |   ],
|   ],
|   |<HistoryStep> BedTool("/tmp/pybedtools.SFmbYc.tmp").merge(),
]
```

```

|         |         parent tag: omfrkwjp,
|         |         result tag: zlwqblvk
|     ],
|
|<HistoryStep> BedTool("/tmp/pybedtools.wlBiMo.tmp").subtract(
|         "/usr/local/lib/python2.6/dist-packages/pybedtools/test/data/b.bed",
|         ),
|         parent tag: zlwqblvk,
|         result tag: reztxhen
|
]
```

Those first three history steps correspond to `c`, `d`, and `e` respectively, as we can see by comparing the code snippet above with the commands in each history step. In other words, `e` can be described by the sequence of 3 commands in the first three history steps. In fact, if we checked `e.history`, we'd see exactly those same 3 steps.

When `f` was created above, it operated both on `e`, which had its own history, as well as `b` – note the nesting of the list. You can do arbitrarily complex “genome algebra” operations, and the history of the `BEDTools` will keep track of this. It may not be useful in every situation, but the ability to backtrack and have a record of what you’ve done can sometimes be helpful.

2.3.12 Deleting temp files specific to a single `BedTool`

You can delete temp files that have been created over the history of a `BedTool` with `BedTool.delete_temporary_history()`. This method will inspect the history, figure out which items point to files in the temp dir (which you can see with `get_tempdir()`), and prompt you for their deletion:

```

>>> f.delete_temporary_history()
Delete these files?
    /tmp/pybedtools..BgULVj.tmp
    /tmp/pybedtools.SFmbYc.tmp
    /tmp/pybedtools.wlBiMo.tmp
(y/N) y
```

Note that the file that `f` points to is left alone. To clarify, the `BedTool.delete_temporary_history()` will only delete temp files that match the pattern `<TEMP_DIR>/pybedtools.*.tmp` from the history of `f`, up to but not including the file for `f` itself. Any `BedTool` instances that do not match the pattern are left alone. Use the kwarg `ask=False` to disable the prompt.

2.4 Topical Documentation

This section contains additional documentation not covered in the tutorial.

2.4.1 Design principles

Hopefully, understanding (or just being aware of) these design principles will help in getting the most out of `pybedtools` and working efficiently.

Principle 1: Temporary files are created automatically

Using `BedTool` instances typically has the side effect of creating temporary files on disk. Even when using the iterator protocol of `BedTool` objects, temporary files may be created in order to run `BEDTools` programs (see [Iterators](#) for more on this latter topic).

Let's illustrate some of the design principles behind `pybedtools` by merging features in a `.bed` that are 100 bp or less apart (`d=100`) in a strand-specific way (`s=True`):

```
>>> from pybedtools import BedTool
>>> import pybedtools
>>> a = BedTool(pybedtools.example_filename('a.bed'))
>>> merged_a = a.merge(d=100, s=True)
```

Now `merged_a` is a `BedTool` instance that contains the results of the merge.

`BedTool` objects must always point to a file on disk. So in the example above, `merged_a` is a `BedTool`, but what file does it point to? You can always check the `BedTool.fn` attribute to find out:

```
>>> # what file does 'merged_a' point to?
>>> merged_a.fn
'/tmp/pybedtools.MPPp5f.tmp'
```

Note that the specific filename will be different for you since it is a randomly chosen name (handled by Python's `tempfile` module). This shows one important aspect of `pybedtools`: every operation results in a new temporary file. Temporary files are stored in `/tmp` by default, and have the form `/tmp/pybedtools.*.tmp`.

When you are done using the `pybedtools` module, make sure to clean up all the temp files created with:

```
>>> # Don't do this yet if you're following the tutorial!
>>> pybedtools.cleanup()
```

If you forget to do this, from the command line you can always do a:

```
rm /tmp/pybedtools.*.tmp
```

to clean everything up. Alternatively, in this session or another session you can use:

```
>>> pybedtools.cleanup(remove_all=True)
```

to remove all files that match the pattern `<tempdir>/pybedtools.*.tmp` where `<tempdir>` is the current value of `pybedtools.get_tempdir()`.

If you need to specify a different directory than that used by default by Python's `tempdir` module, then you can set it with:

```
>>> pybedtools.set_tempdir('/scratch')
```

You'll need write permissions to this directory, and it needs to already exist. All temp files will then be written to that directory, until the `tempdir` is changed again.

Principle 2: Names and arguments are as similar as possible to BEDTools

As much as possible, `BEDTools` programs and `BedTool` methods share the same names and arguments.

Returning again to this example:

```
>>> merged_a = a.merge(d=100, s=True)
```

This demonstrates that the `BedTool` methods that wrap `BEDTools` programs do the same thing and take the exact same arguments as the `BEDTools` program. Here we can pass `d=100` and `s=True` only because the underlying `BEDTools` program, `mergeBed`, can accept these arguments. Need to know what arguments `mergeBed` can take? See the docs for `BedTool.merge()`; for more on this see *Principle 7: Check the help*.

In general, remove the “Bed” from the end of the `BEDTools` program to get the corresponding `BedTool` method. So there's a `BedTool.subtract()` method for `subtractBed`, a `BedTool.intersect()` method for `intersectBed`, and so on.

Principle 3: Indifference to BEDTools version

Since `BedTool` methods just wrap `BEDTools` programs, they are as up-to-date as the version of `BEDTools` you have installed on disk. If you are using a cutting-edge version of `BEDTools` that has some hypothetical argument `-z` for `intersectBed`, then you can use `a.intersectBed(z=True)`.

`pybedtools` will also raise an exception if you try to use a method that relies on a more recent version of `BEDTools` than you have installed.

Principle 4: Sensible default args

If we were running the `mergeBed` program from the command line, we would have to specify the input file with the `mergeBed -i` option.

`pybedtools` assumes that if we're calling the `merge()` method on the `BedTool`, `a`, we want to operate on the bed file that `a` points to.

In general, `BEDTools` programs that accept a single BED file as input (by convention typically specified with the `-i` option) the default behavior for `pybedtools` is to use the `BedTool`'s file (indicated in the `BedTool.fn` attribute) as input.

We can still pass a file using the `i` keyword argument if we wanted to be absolutely explicit. In fact, the following two versions produce the same output:

```
>>> # The default is to use existing file for input -- no need
>>> # to specify "i" . . .
>>> result1 = a.merge(d=100, s=True)

>>> # . . . but you can always be explicit if you'd like
>>> result2 = a.merge(i=a.fn, d=100, s=True)

>>> # Confirm that the output is identical
>>> result1 == result2
True
```

Methods that have this type of default behavior are indicated by the following text in their docstring:

```
.. note::

    For convenience, the file this BedTool object points to is passed as "-i"
```

There are some `BEDTools` programs that accept two BED files as input, like `intersectBed` where the first file is specified with `-a` and the second file with `-b`. The default behavior for `pybedtools` is to consider the `BedTool`'s file as `-a` and the first non-keyword argument to the method as `-b`, like this:

```
>>> b = pybedtools.example_bedtool('b.bed')
>>> result3 = a.intersect(b)
```

This is exactly the same as passing the `a` and `b` arguments explicitly:

```
>>> result4 = a.intersect(a=a.fn, b=b.fn)
>>> result3 == result4
True
```

Furthermore, the first non-keyword argument used as `-b` can either be a filename *or* another `BedTool` object; that is, these commands also do the same thing:

```
>>> result5 = a.intersect(b=b.fn)
>>> result6 = a.intersect(b=b)
```

```
>>> str(result5) == str(result6)
True
```

Methods that accept either a filename or another `BedTool` instance as their first non-keyword argument are indicated by the following text in their docstring:

```
.. note::

    This method accepts either a BedTool or a file name as the first
    unnamed argument
```

Principal 5: Other arguments have no defaults

Only the `BEDTools` arguments that refer to BED (or other interval) files have defaults. In the current version of `BEDTools`, this means only the `-i`, `-a`, and `-b` arguments have defaults. All others have no defaults specified by `pybedtools`; they pass the buck to `BEDTools` programs. This means if you do not specify the `d` kwarg when calling `BedTool.merge()`, then it will use whatever the installed version of `BEDTools` uses for `-d` (currently, `mergeBed`'s default for `-d` is 0).

`-d` is an option to `BEDTools mergeBed` that accepts a value, while `-s` is an option that acts as a switch. In `pybedtools`, simply pass a value (integer, float, whatever) for value-type options like `-d`, and boolean values (`True` or `False`) for the switch-type options like `-s`.

Here's another example using both types of keyword arguments; the `BedTool` object `b` (or it could be a string filename too) is implicitly passed to `intersectBed` as `-b` (see *Principle 4: Sensible default args* above):

```
>>> a.intersect(b, v=True, f=0.5)
```

Again, any option that can be passed to a `BEDTools` program can be passed to the corresponding `BedTool` method.

Principle 6: Chaining together commands

Most methods return new `BedTool` objects, allowing you to chain things together just like piping commands together on the command line. To give you a flavor of this, here is how you would get the merged regions of features shared between `a.bed` (as referred to by the `BedTool` `a` we made previously) and `b.bed`: (as referred to by the `BedTool` `b`):

```
>>> a.intersect(b).merge().saveas('shared_merged.bed')
<BedTool(shared_merged.bed)>
```

This is equivalent to the following `BEDTools` commands:

```
intersectBed -a a.bed -b b.bed | merge -i stdin > shared_merged.bed
```

Methods that return a new `BedTool` instance are indicated with the following text in their docstring:

```
.. note::

    This method returns a new BedTool instance
```

Principle 7: Check the help

If you're unsure of whether a method uses a default, or if you want to read about what options an underlying `BEDTools` program accepts, check the help. Each `pyBedTool` method that wraps a `BEDTools` program also wraps the `BEDTools` program help string. There are often examples of how to use a method in the docstring as well. The documentation is also run through doctests, so the code you read here is guaranteed to work and be up-to-date.

2.4.2 Creating a BedTool

To create a `BedTool`, first you need to import the `pybedtools` module. For these examples, I’m assuming you have already done the following:

```
>>> import pybedtools
>>> from pybedtools import BedTool
```

Next, you need a BED file to work with. If you already have one, then great – move on to the next section. If not, `pybedtools` comes with some example bed files used for testing. You can take a look at the list of example files that ship with `pybedtools` with the `list_example_files()` function:

```
>>> # list the example bed files
>>> pybedtools.list_example_files()
['a.bed', 'b.bed', 'c.gff', 'd.gff', 'dm3-chr2L-5M-invalid.gff.gz', 'dm3-chr2L-5M.gff.gz', 'dmel-all-
```

Once you decide on a file to use, feed the your choice to the `example_filename()` function to get the full path:

```
>>> # get the full path to an example bed file
>>> bedfn = pybedtools.example_filename('a.bed')
```

The full path of `bedfn` will depend on your installation (this is similar to the `data()` function in `R`, if you’re familiar with that).

Now that you have a filename – either one of the example files or your own, you create a new `BedTool` simply by pointing it to that filename:

```
>>> # create a new BedTool from the example bed file
>>> myBedTool = BedTool(bedfn)
```

Alternatively, you can construct BED files from scratch by using the `from_string` keyword argument. However, all spaces will be converted to tabs using this method, so you’ll have to be careful if you add “name” columns. This can be useful if you want to create *de novo* BED files on the fly:

```
>>> # an "inline" example:
>>> fromscratch1 = pybedtools.BedTool('chrX 1 100', from_string=True)
>>> print fromscratch1
chrX    1    100
```

```
>>> # using a longer string to make a bed file. Note that
>>> # newlines don't matter, and one or more consecutive
>>> # spaces will be converted to a tab character.
>>> larger_string = """
... chrX 1    100    feature1  0  +
... chrX 50   350    feature2  0  -
... chr2 5000 10000  another_feature 0  +
... """
```

```
>>> fromscratch2 = BedTool(larger_string, from_string=True)
>>> print fromscratch2
chrX    1    100 feature1    0    +
chrX    50   350 feature2    0    -
chr2    5000  10000  another_feature 0    +
```

Of course, you’ll usually be using your own bed files that have some biological importance for your work that are saved in places convenient for you, for example:

```
>>> a = BedTool('/data/sample1/peaks.bed')
```

2.4.3 Using BedTool objects as iterators/generators

Typically, `BedTool` objects are used somewhat like handles to individual files on disk that contain BED lines. To save disk space, `BedTool` objects also have the ability to “stream”, much like piping in Unix. That is, the data are created only one line at a time in memory, instead of either creating a list of all data in memory or writing all data to disk.

Warning: You’ll need to be careful when using `BedTool` objects as generators, since any operation that reads all the features of a `BedTool` will consume the iterable.

To get a streaming `BedTool`, use the `stream=True` kwarg. This `BedTool` will act a little differently from a standard, file-based `BedTool`.

```
>>> a = pybedtools.example_bedtool('a.bed')
>>> b = pybedtools.example_bedtool('b.bed')
>>> c = a.intersect(b, stream=True)

>>> # checking the length consumes the iterator
>>> len(c)
3

>>> # nothing left, so checking length again returns 0
>>> len(c)
0
```

In some cases, a stream may be “rendered” to a temp file. This is because BEDTools programs can only accept one input file as `stdin`. This is typically the first input (`-i` or `-a`), while the other input (`-b`) must be a file. Consider this example, where the second intersection needs to convert the streaming `BedTool` to a file before sending to `intersectBed`:

```
>>> a = pybedtools.example_bedtool('a.bed')
>>> b = pybedtools.example_bedtool('b.bed')

>>> # first we set up a streaming BedTool:
>>> c = a.intersect(b, stream=True)

>>> # But supplying a streaming BedTool as the first unnamed argument
>>> # means it is being passed as -b to intersectBed, and so must be a file.
>>> # In this case, 'c' is rendered to a tempfile before being passed.
>>> d = a.intersect(c, stream=True)
```

Creating a `BedTool` from an iterable

You can create a `BedTool` on the fly from a generator or iterator – in fact, this is what the `BedTool.filter()` method does for you:

```
>>> a = pybedtools.example_bedtool('a.bed')
>>> print a
chr1      1      100      feature1      0      +
chr1     100     200      feature2      0      +
chr1     150     500      feature3      0      -
chr1     900     950      feature4      0      +

>>> b = pybedtools.BedTool(f for f in a if f.start > 200)
```

```
>>> # this is the same as using filter:
>>> c = a.filter(lambda x: x.start > 200)
```

We need to “render” these BedTools to string before we can check equality – consuming them both – since they are both iterables for which == is not defined:

```
>>> b == c
Traceback (most recent call last):
...
NotImplementedError: Testing equality only supported for BedTools that point to a file

>>> str(b) == str(c)
True
```

2.4.4 Working with BAM files

Some BEDTools programs support BAM files as input; for example `intersectBed`, `windowBed`, and others accept a `-abam` argument instead of `-a` for the first input file.

This section describes the workflow for working with BAM files within `pybedtools`.

As an example, let’s intersect a BAM file of reads with annotations using files that ship with `pybedtools`. First, we create the BedTool objects:

```
>>> a = pybedtools.example_bedtool('x.bam')
>>> b = pybedtools.example_bedtool('dm3-chr2L-5M.gff.gz')
```

If `a` referred to a BED file like `a.bed`, we could just do `a.intersect(b)` because `a.bed` would be implicitly passed as `-a` and the gzipped GFF file would be passed as `-b`. In order to use a BAM file, however, we need to explicitly specify an `abam` kwarg. In addition, since Python doesn’t allow non-keyword arguments after keyword arguments, we need to explicitly specify a `b` kwarg.

This should be much clearer with a simple example:

```
>>> c = a.intersect(abam=a.fn, b=b)
```

Now `c` points to a new BAM file on disk. Keep in mind that there is not yet iterable BAM support in `pybedtools`, so things like `c.count()` or iterating over `c` with a `for feature in c: ...` (i.e., as described in [Using BedTool objects as iterators/generators](#) for BED files) won’t work. For now, consider using a package like `HTSeq` for access to individual reads in BAM format.

Alternatively, we can specify the `bed=True` kwarg to convert the intersected BAM results to BED format, and use those like a normal BED file:

```
>>> d = a.intersect(abam=a.fn, b=b, bed=True)
```

The resulting BedTool `d` refers to a BED file and can be used like any other:

```
>>> d.count()
341324

>>> print iter(d).next()
chr2L      9329      9365      HWUSI-NAME:2:69:512:1017#0      3      -
```

2.4.5 Specifying genomes

This section illustrates the use of genome files for use with BEDTools programs that need to know chromosome limits to prevent out-of-range coordinates.

Using BEDTools programs like `slopBed` or `shuffleBed` from the command line requires “genome” or “chrom-sizes” files. `pybedtools` comes with common genome assemblies already set up as a dictionary with chromosomes as keys and zero-based (start, stop) tuples as values:

```
>>> from pybedtools import genome_registry
>>> genome_registry.dm3['chr2L']
(0, 23011544)
```

The rules for specifying a genome for methods that require a genome are as follows (use whatever is most convenient):

- Use `g` to specify either a filename or a dictionary
- Use `genome` to specify either an assembly name or a dictionary

Below are examples of each.

As a file

This is the typical way of using BEDTools programs, by specifying an existing genome file with `g`:

```
>>> a = pybedtools.example_bedtool('a.bed')
>>> b = a.slop(b=100, g='hg19.genome')
```

As a string

This is probably the most convenient way of specifying a genome. If the genome exists in the genome registry it will be used directly; otherwise it will automatically be downloaded from UCSC. You must use the `genome` kwarg for this; if you use `g` a string will be interpreted as a filename:

```
>>> c = a.slop(b=100, genome='hg19')
```

As a dictionary

This is a good way of providing custom coordinates; either `g` or `genome` will accept a dictionary:

```
>>> d = a.slop(b=100, g={'chr1': (1, 10000)})
>>> e = a.slop(b=100, genome={'chr1': (1, 10000)})
```

Make sure that all these different methods return the same results

```
>>> b == c == d == e
True
```

Converting to a file

Since **BEDTools** programs operate on files, the fastest choice will be to use an existing file. While the time to convert a dictionary to a file is extremely small, over 1000's of files (e.g., for Monte Carlo simulations), the time may add up. The function `pybedtools.chromsizes_to_file()` will create a file from a dictionary or string:

```
>>> # with no filename specified, a tempfile will be created
>>> pybedtools.chromsizes_to_file(pybedtools.chromsizes('dm3'), 'dm3.genome')
'dm3.genome'
>>> print open('dm3.genome').read()
chr2L      23011544
chr2LHet    368872
chr2R      21146708
chr2RHet    3288761
chr3L      24543557
chr3LHet    2555491
chr3R      27905053
chr3RHet    2517507
chr4       1351857
chrM       19517
chrU       10049037
chrUextra   29004656
chrX       22422827
chrXHet    204112
chrYHet    347038
```

2.4.6 Randomization

`pybedtools` provides some basic functionality for assigning some significance value to the overlap between two BEDfiles.

The strategy is to randomly shuffle a file many times, each time doing an intersection with another file of interest and counting the number of intersections. Upon doing this many times, an empirical distribution is constructed, and the number of intersections between the original, un-shuffled file is compared to this empirical distribution to obtain a p-value, or compared to the median of the distribution to get a score.

There are two methods, `pybedtools.BedTool.randomintersection()` which does the brute force randomizations, and `BedTool.randomstats()` which compiles and reports the results from the former method.

Example workflow

As a somewhat trivial example, we'll intersect the example `a.bed` with `b.bed`, taking care to set some options that will let it run in a deterministic way so that these tests will run.

We will be shuffling `a.bed`, so we'll need to specify the limits of its chromosomes with `BedTool.set_chromsizes()`. Here, we set it to an artificially small chromosome size so that we can get some meaningful results in reasonable time. In practice, you would either supply your own dictionary or use a string assembly name (e.g., `'hg19'`, `'mm9'`, `'dm3'`, etc). The genome-handling code will find the chromsizes we've set, so there's no need to tell `shuffleBed` which genome file to use each time.

```
>>> chromsizes = {'chr1': (0, 1000)}
>>> a = pybedtools.example_bedtool('a.bed').set_chromsizes(chromsizes)
>>> b = pybedtools.example_bedtool('b.bed')
```

We have the option of specifying what kwargs to provide `BedTool.shuffle()` and `BedTool.intersect()`, which will be called each iteration. In this example, we'll tell `shuffleBed` to only shuffle within the chromosome just to illustrate the kwargs passing. We also need to specify how many iterations to perform. In practice, 1000 or 10000 are good numbers, but for the sake of this example we'll only do 100.

Last, setting `debug=True` means that the random seed will be set in a predictable manner so that we'll always get the same results for testing. In practice, make sure you use `debug=False` (the default) to ensure random results.

```
>>> results = a.randomintersection(b, iterations=100, shuffle_kwargs={'chrom': True}, debug=True)
```

results is a generator of intersection counts where each number is the number of times the shuffled a intersected with b. We need to convert it to a list in order to look at it:

```
>>> results = list(results)
>>> len(results)
100

>>> print results[:10]
[1, 1, 2, 2, 1, 2, 1, 0, 2, 3]
```

Running thousands of iterations on files with many features will of course result in more complex results. We could then take these results and plot them in matplotlib, or get some statistics on them.

The method `BedTool.randomstats()` does this for us, but requires NumPy and SciPy to be installed. This method also calls `BedTool.randomintersection()` for us, returning the summarized results in a dictionary.

`BedTool.randomstats()` takes the same arguments as `BedTool.randomintersection()`:

```
>>> results_dict = a.randomstats(b, iterations=100, shuffle_kwargs={'chrom': True}, debug=True)
```

The keys to this results dictionary are as follows (some are redundant, I've found these keys useful for writing out to file):

- iterations** the number of iterations we specified
- actual** the number of intersections between then un-shuffled a and b
- file_a** the filename of a
- file_b** the filename of b
- <file_a>** the key is actually the filename of a, and the value is the number of features in a
- <file_b>** the key is actually the filename of b and the value is the number of features in b
- self** number of features in a (or "self"; same value as for <file_a>)
- other** number of features in b (or "other"; same value as for <file_b>)
- frac randomized above actual** fraction of iterations that had counts above the actual count
- frac randomized below actual** fraction of iterations that had counts below the actual count
- median randomized** the median of the distribution of randomized intersections
- normalized** the actual count divided by the median; can be considered as a score
- percentile** the percentile of actual within the distribution of randomized intersections; can be considered an empirical p-value
- upper 97.5th** the 97.5th percentile of the randomized distribution
- lower 2.5th** the 2.5th percentile of the randomized distribution

For example:

```
>>> keys = ['self', 'other', 'actual', 'median randomized', 'normalized', 'percentile']
>>> for key in keys:
...     print '%s: %s' % (key, results_dict[key])
self: 4
other: 2
actual: 3
median randomized: 2.0
```



```
normalized: 1.5
percentile: 92.0
```

Contributions toward improving this code or implementing other methods of statistical testing are very welcome!

2.4.7 Wrapping new tools

This section serves as a reference for wrapping new tools as they are added to BEDTools.

Let's assume we would like to wrap a new program, appropriately named `newProgramBed`. Its signature from the command line is `newProgramBed -a <infile> -b <other file> [options]`, and it accepts `-a stdin` to indicate data is being piped to it.

Method name

Generally, I've tried to keep method names as similar as possible to BEDTools programs while still being PEP8-compliant. The trailing 'Bed' is usually removed from the program name. So here the name would probably be `new_program`.

Method signature (args and kwargs)

If a BEDTools program accepts a `-b` argument (which is the case for this example) then the signature and first line should look like this:

```
def new_program(self, b=None, **kwargs):
    kwargs['b'] = b
    ...
```

That is, we specify a kwarg for `b` in the signature so that we have the option calling this method either with a regular arg or a kwarg, and then we ensure that `b` makes it into the `**kwargs` dictionary that will be passed around later.

This is to allow another BedTool (or file, or stream) be passed as the first non-keyword argument; otherwise, the user would always have to do:

```
a.intersect(b=b)
```

instead of:

```
a.intersect(b)
```

It's a minor thing, but it's convenient.

If the program you're wrapping doesn't accept another BED-like file as `-b`, then the method should only accept `self` and `**kwargs`:

```
def new_program(self, **kwargs):
    pass
```

Setting the default

Generally, the next thing to do is to set the default, or implicit kwarg. This is generally the `-i` or `-a` kwarg.

This is done by checking to see if the implicit kwarg was specified already (which allows the user to override the implicit assumption) and if not, add it to the `kwargs` dict.

For example:

```
if 'a' not in kwargs:
    kwargs['a'] = self.fn
```

For BEDTools programs that optionally take an `-abam` argument, we need to check to make sure that wasn't specified either:

```
if ('abam' not in kwargs) and ('a' not in kwargs):
    kwargs['a'] = self.fn
```

So to continue the example, our method now looks like this:

```
def new_program(self, b=None, **kwargs):

    if 'a' not in kwargs:
        kwargs['a'] = self.fn
```

Allowing input streams

BEDTools programs that can accept `stdin` as their first input need to be registered in the `BedTool.handle_kwargs()` method, in the `implicit_instream1` dictionary.

This dictionary specifies which keyword argument to look for an object that could be used as the first input – this could be a filename, an iterator of features, or an open file. Depending on what it finds there, it will call the BEDTools program appropriately.

For example, if `BedTool.handle_kwargs()` finds an open file in `kwargs['a']` for `intersectBed`, then it will tell `intersectBed` that `-a` should be `stdin` and the open file will eventually be passed to the `subprocess.Popen` call. But if `kwargs['a']` is a filename, then it will just tell `intersectBed` that `-a` should be the filename.

If there is a second argument that has the potential to be a stream (like the `b` kwarg for `BedTool.intersect()`, which can be a `BedTool`, filename, `IntervalIterator`, or stream), then this kwarg should be added to the `implicit_instream2` dictionary.

This second dictionary specifies which kwargs to check to see if they contain an iterable. If so, then it “renders” the iterable to a temp file and will pass that filename to the BEDTools program.

In the example case of `BedTool.intersect()`, if `b` is a stream or an iterable, the `handle_kwargs` method will convert that stream or iterable to a tempfile (say, `tmp001`) and then will eventually send `intersectBed` the option `-b tmp001`.

The `BedTool.handle_kwargs()` method also decides whether to create a new temp file (if `stream=False`) or not, and also converts kwargs like `a` to `-a`. Finally, it creates a list of commands ready for `call_bedtools()` to run.

To illustrate, we add `a` to the `implicit_instream1` dict:

```
implicit_instream1 = {'intersectBed': 'a',
                      'subtractBed': 'a',
                      'closestBed': 'a',
                      'windowBed': 'a',
                      'slopBed': 'i',
                      'mergeBed': 'i',
                      'sortBed': 'i',
                      'bed12ToBed6': 'i',
                      'shuffleBed': 'i',
                      'annotateBed': 'i',
                      'flankBed': 'i',
                      'fastaFromBed': 'bed',
```

```

        'maskFastaFromBed': 'bed',
        'coverageBed': 'a',
        'newthingBed': 'a', # here's the new wrapped program
    }

```

And back in the method body, we call `BedTool.handle_kwargs()` so that our method now looks like this:

```

def new_program(self, b=None, **kwargs):

    if 'a' not in kwargs:
        kwargs['a'] = self.fn

    cmds, tmp, stdin = self.handle_kwargs(prog='newthingBed', **kwargs)

```

Call BEDTools, and return result

So now we have `cmds` (the list of commands ready to be called by `subprocess.Popen`), `tmp` (the new tempfile ready to accept results, or `None` if we will be streaming the output), and `stdin` (`None` if it's a filename, or the open file that will be send to `subprocess.stdin`). These are passed to the `call_bedtools()` function, which does all the subprocess business and returns a “stream”, which could be any of the things that a `BedTool` can accept (filename, open file, `IntervalIterator`). Finally, we return a new `BedTool` made from this stream.

Now our method looks like this:

```

def new_program(self, b=None, **kwargs):

    if 'a' not in kwargs:
        kwargs['a'] = self.fn

    cmds, tmp, stdin = self.handle_kwargs(prog='newthingBed', **kwargs)
    stream = call_bedtools(cmds, tmp, stdin=stdin)
    return BedTool(stream)

```

Add decorators

Some decorators are used to add text to the method's docstring, like `@implicit` and `@file_or_bedtool`, and `@returns_bedtool`. More useful is `@help`, which adds the full text of the BEDTools program to the end of the docstring, so all information is available from the Python interpreter (especially useful when using IPython). The `@log_to_history` decorator will register the calling of this method in the `BedTool`'s history.

The final wrapped method, with all the decorators to add relevant text to the docstring, then is simply:

```

@returns_bedtool()
@file_or_bedtool()
@help('newProgramBed')
@log_to_history
@implicit('-a')
def new_program(self, b=None, **kwargs):

    if 'a' not in kwargs:
        kwargs['a'] = self.fn

    cmds, tmp, stdin = self.handle_kwargs(prog='newthingBed', **kwargs)
    stream = call_bedtools(cmds, tmp, stdin=stdin)
    return BedTool(stream)

```

Write tests!

The only way to know for sure if your new wrapped method works is to write good tests for it. This can either be done in the docstring or in the test suite. See the source for how this is done; also check out the `test.sh` script in the top level of the repository.

Send a pull request

If you've made something that you think would be useful to others, please send a github pull request so that your newly created and tested code can be distributed to others. Any contributions are much appreciated.

2.5 pybedtools Reference

This section is the module reference documentation, and includes the full docstrings for methods and functions in `pybedtools`. It is separated into *pybedtools module-level functions*, *BedTool methods that wrap BEDTools programs*, and *BedTool methods unique to pybedtools*.

Contents

- `pybedtools` Reference
 - `pybedtools` module-level functions
 - * Functions for working with example files
 - * Functions for specifying genome assemblies
 - * Setup
 - * Utilities
 - `BedTool` methods that wrap `BEDTools` programs
 - * “Genome algebra” methods
 - `intersect()` (wraps “`intersectBed`”)
 - `merge()` (wraps “`mergeBed`”)
 - `subtract()` (wraps “`subtractBed`”)
 - `closest()` (wraps “`closestBed`”)
 - `window()` (wraps “`windowBed`”)
 - `sort()` (wraps “`sortBed`”)
 - `slop()` (wraps “`slopBed`”)
 - `flank()` (wraps “`flankBed`”)
 - `shuffle()` (wraps “`shuffleBed`”)
 - `annotate()` (wraps “`annotateBed`”)
 - `coverage()` (wraps “`coverageBed`”)
 - `genome_coverage()` (wraps “`genomeCoverageBed`”)
 - * Methods for converting between formats
 - `bed6()` (wraps “`Bed12To6`”)
 - * Methods for working with sequences
 - `sequence()` (wraps “`fastaFromBed`”)
 - `mask_fasta()` (wraps “`maskFastaFromBed`”)
 - `BedTool` methods unique to `pybedtools`
 - * Introspection
 - `count()`
 - `print_sequence()`
 - `field_count()`
 - * Saving
 - `saveas()`
 - `save_seqs()`
 - * Utilities
 - `with_attrs()`
 - `cat()`
 - `total_coverage()`
 - `delete_temporary_history()`
 - * Feature-by-feature operations
 - `each()`
 - `filter()`
 - `cut()`
 - `features()`
 - * Randomization helpers
 - `randomintersection()`
 - `randomstats()`

2.5.1 pybedtools module-level functions

Functions for working with example files

`pybedtools.example_bedtool(fn)`

Return a bedtool using a bed file from the pybedtools examples directory. Use `list_example_files()` to see a list of files that are included.

`pybedtools.example_filename(fn)`

Return a bed file from the pybedtools examples directory. Use `list_example_files()` to see a list of files that are included.

`pybedtools.list_example_files()`

Returns a list of files in the examples dir. Choose one and pass it to `example_file_fnl()` to get the full path to an examplefile.

Example usage:

```
>>> choices = list_example_files()
>>> assert 'a.bed' in choices
>>> bedfn = example_filename('a.bed')
>>> mybedtool = BedTool(bedfn)
```

`pybedtools.data_dir()`

Returns the data directory that contains example files for tests and documentation.

Functions for specifying genome assemblies

`pybedtools.chromsizes(genome)`

Looks for a *genome* already included in the genome registry; if not found then it looks it up on UCSC. Returns the dictionary of chromsize tuples where each tuple has (start,stop).

Chromsizes are described as (start, stop) tuples to allow randomization within specified regions; e. g., you can make a chromsizes dictionary that represents the extent of a tiling array.

Example usage:

```
>>> dm3_chromsizes = chromsizes('dm3')
>>> for i in sorted(dm3_chromsizes.items()):
...     print i
('chr2L', (0, 23011544))
('chr2LHet', (0, 368872))
('chr2R', (0, 21146708))
('chr2RHet', (0, 3288761))
('chr3L', (0, 24543557))
('chr3LHet', (0, 2555491))
('chr3R', (0, 27905053))
('chr3RHet', (0, 2517507))
('chr4', (0, 1351857))
('chrM', (0, 19517))
('chrU', (0, 10049037))
('chrUextra', (0, 29004656))
('chrX', (0, 22422827))
('chrXHet', (0, 204112))
('chrYHet', (0, 347038))
```

`pybedtools.chromsizes_to_file(chromsizes,fn=None)`

Converts a *chromsizes* dictionary to a file. If *fn* is None, then a tempfile is created (which can be deleted with `pybedtools.cleanup()`).

Returns the filename.

`pybedtools.get_chromsizes_from_ucsc(genome, saveas=None, mysql='mysql', timeout=None)`
Download chrom size info for *genome* from UCSC and returns the dictionary.

If you need the file, then specify a filename with *saveas* (the dictionary will still be returned as well).

If *mysql* is not on your path, specify where to find it with *mysql=<path to mysql executable>*.

timeout is how long to wait for a response; mostly used for testing.

Example usage:

```
>>> dm3_chromsizes = get_chromsizes_from_ucsc('dm3')
>>> for i in sorted(dm3_chromsizes.items()):
...     print i
('chr2L', (0, 23011544))
('chr2LHet', (0, 368872))
('chr2R', (0, 21146708))
('chr2RHet', (0, 3288761))
('chr3L', (0, 24543557))
('chr3LHet', (0, 2555491))
('chr3R', (0, 27905053))
('chr3RHet', (0, 2517507))
('chr4', (0, 1351857))
('chrM', (0, 19517))
('chrU', (0, 10049037))
('chrUextra', (0, 29004656))
('chrX', (0, 22422827))
('chrXHet', (0, 204112))
('chrYHet', (0, 347038))
```

Setup

`pybedtools.set_tempdir(tempdir)`

Sets the directory for temp files. Useful for clusters that use a /scratch partition rather than a /tmp dir. Convenience function to simply set `tempfile.tempdir`.

`pybedtools.get_tempdir()`

Gets the current tempdir for the module.

`pybedtools.set_bedtools_path(path='')`

If BEDTools is not available on your system path, specify the path to the dir containing the BEDTools executables (`intersectBed`, `subtractBed`, etc) with this function.

To reset and use the default system path, call this function with no arguments or use `path=''`.

Utilities

`pybedtools.cleanup(verbose=True, remove_all=False)`

Deletes all temporary files in the `BedTool.TEMPFILES` class variable.

If *verbose*, reports what it's doing

If *remove_all*, then ALL files matching “`pybedtools*.tmp`” in the temp dir will be deleted.

`pybedtools.IntervalIterator(stream)`

Given an open file handle, yield the Intervals.

`pybedtools.find_tagged(tag)`

Returns the bedtool object with tagged with *tag*. Useful for tracking down bedtools you made previously.

2.5.2 BedTool methods that wrap BEDTools programs

“Genome algebra” methods

`intersect()` (wraps “`intersectBed`”)

`BedTool.intersect(*args, **kwargs)`

pybedtools help:

Intersect with another BED file. If you want to use BAM as input, you need to specify *abam*=*'filename.bam'*. Returns a new BedTool object.

Example usage:

Create new BedTool object

```
>>> a = pybedtools.example_bedtool('a.bed')
```

Get overlaps with *b.bed*:

```
>>> b = pybedtools.example_bedtool('b.bed')
>>> overlaps = a.intersect(b)
```

Use *v=True* to get the inverse – those unique to “*a.bed*”:

```
>>> unique_to_a = a.intersect(b, v=True)
```

Note: This method returns a new bedtool instance

Note: For convenience, the file this bedtool object points to is passed as “*-a*”

Note: This method accepts either a bedtool or a file name as the first unnamed argument

Original BEDtools program help:

Program: `intersectBed` (v2.12.0) Author: Aaron Quinlan (aaronquinlan@gmail.com) Summary: Report overlaps between two feature files.

Usage: `intersectBed [OPTIONS] -a <bed/gff/vcf> -b <bed/gff/vcf>`

Options:

-abam	The A input file is in BAM format. Output will be BAM as well.
-ubam	Write uncompressed BAM output. Default is to write compressed BAM.
-bed	When using BAM input (<i>-abam</i>), write output as BED. The default is to write output in BAM when using <i>-abam</i> .
-wa	Write the original entry in A for each overlap.

-wb	Write the original entry in B for each overlap. - Useful for knowing what A overlaps. Restricted by -f and -r.
-wo	Write the original A and B entries plus the number of base pairs of overlap between the two features. - Overlaps restricted by -f and -r. Only A features with overlap are reported.
-wao	Write the original A and B entries plus the number of base pairs of overlap between the two features. - Overlapping features restricted by -f and -r. However, A features w/o overlap are also reported with a NULL B feature and overlap = 0.
-u	Write the original A entry once if any overlaps found in B. - In other words, just report the fact ≥ 1 hit was found. - Overlaps restricted by -f and -r.
-c	For each entry in A, report the number of overlaps with B. - Reports 0 for A entries that have no overlap with B. - Overlaps restricted by -f and -r.
-v	Only report those entries in A that have no overlaps with B. - Similar to “grep -v” (an homage).
-f	Minimum overlap required as a fraction of A. - Default is 1E-9 (i.e., 1bp). - FLOAT (e.g. 0.50)
-r	Require that the fraction overlap be reciprocal for A and B. - In other words, if -f is 0.90 and -r is used, this requires that B overlap 90% of A and A also overlaps 90% of B.
-s	Force strandedness. That is, only report hits in B that overlap A on the same strand. - By default, overlaps are reported without respect to strand.
-split	Treat “split” BAM or BED12 entries as distinct BED intervals.

`merge()` (wraps “mergeBed”)

`BedTool.merge(*args, **kwargs)`
pybedtools help:

Merge overlapping features together. Returns a new BedTool object.

Example usage:

```
>>> a = pybedtools.example_bedtool('a.bed')
```

Merge:

```
>>> c = a.merge()
```

Allow merging of features 500 bp apart:

```
>>> c = a.merge(d=500)
```

Report number of merged features:

```
>>> c = a.merge(n=True)
```

Report names of merged features:

```
>>> c = a.merge(nms=True)
```

Note: This method returns a new bedtool instance

Note: For convenience, the file this bedtool object points to is passed as “-i”

Original BEDtools program help:

Program: mergeBed (v2.12.0) Author: Aaron Quinlan (aaronquinlan@gmail.com) Summary: Merges overlapping BED/GFF/VCF entries into a single interval.

Usage: mergeBed [OPTIONS] -i <bed/gff/vcf>

Options:

-s	Force strandedness. That is, only merge features that are the same strand. - By default, merging is done without respect to strand.
-n	Report the number of BED entries that were merged. - Note: “1” is reported if no merging occurred.
-d	Maximum distance between features allowed for features to be merged. - Def. 0. That is, overlapping & book-ended features are merged. - (INTEGER)
-nms	Report the names of the merged features separated by semi-colons.

-scores [STRING] Report the scores of the merged features. Specify one of

the following options for reporting scores: sum, min, max, mean, median, mode, anti-mode, collapse (i.e., print a semicolon-separated list),

subtract () (wraps “subtractBed”)

BedTool.**subtract** (*args, **kwargs)

pybedtools help:

Subtracts from another BED file and returns a new BedTool object.

Example usage:

```
>>> a = pybedtools.example_bedtool('a.bed')
>>> b = pybedtools.example_bedtool('b.bed')
```

Do a “stranded” subtraction:

```
>>> c = a.subtract(b, s=True)
```

Require 50% of features in a to overlap:

```
>>> c = a.subtract(b, f=0.5)
```

Note: This method returns a new bedtool instance

Note: This method accepts either a bedtool or a file name as the first unnamed argument

Original BEDtools program help:

Program: subtractBed (v2.12.0) Author: Aaron Quinlan (aaronquinlan@gmail.com) Summary: Removes the portion(s) of an interval that is overlapped

by another feature(s).

Usage: subtractBed [OPTIONS] -a <bed/gff/vcf> -b <bed/gff/vcf>

Options:

-f	Minimum overlap required as a fraction of A. - Default is 1E-9 (i.e., 1bp). - (FLOAT) (e.g. 0.50)
-s	Force strandedness. That is, only report hits in B that overlap A on the same strand. - By default, overlaps are reported without respect to strand.

closest () (wraps “closestBed”)

BedTool.**closest** (*args, **kwargs)

pybedtools help:

Return a new BedTool object containing closest features in *b*. Note that the resulting file is no longer a valid BED format; use the special “_closest” methods to work with the resulting file.

Example usage:

```
a = BedTool('in.bed')

# get the closest feature in 'other.bed' on the same strand
b = a.closest('other.bed', s=True)
```

Note: This method returns a new bedtool instance

Note: For convenience, the file this bedtool object points to is passed as “-a”

Note: This method accepts either a bedtool or a file name as the first unnamed argument

Original BEDtools program help:

Program: closestBed (v2.12.0) Authors: Aaron Quinlan (aaronquinlan@gmail.com)

Erik Arner, Riken

Summary: For each feature in A, finds the closest feature (upstream or downstream) in B.

Usage: closestBed [OPTIONS] -a <bed/gff/vcf> -b <bed/gff/vcf>

Options:

- s** Force strandedness. That is, find the closest feature in B that overlaps A on the same strand. - By default, overlaps are reported without respect to strand.
- d** In addition to the closest feature in B, report its distance to A as an extra column. - The reported distance for overlapping features will be 0.
- t** How ties for closest feature are handled. This occurs when two features in B have exactly the same overlap with A. By default, all such features in B are reported. Here are all the options: - “all” Report all ties (default). - “first” Report the first tie that occurred in the B file. - “last” Report the last tie that occurred in the B file.

Notes: Reports “none” for chrom and “-1” for all other fields when a feature is not found in B on the same chromosome as the feature in A. E.g. none -1 -1

`window()` (wraps “`windowBed`”)

`BedTool.window(*args, **kwargs)`
pybedtools help:

Intersect with a window.

Example usage:

```
>>> a = pybedtools.example_bedtool('a.bed')
>>> b = pybedtools.example_bedtool('b.bed')
>>> print a.window(b, w=1000)
chr1      1      100    feature1      0      +      chr1      155      200    feat
chr1      1      100    feature1      0      +      chr1      800      901    feat
chr1     100     200    feature2      0      +      chr1      155      200    feat
chr1     100     200    feature2      0      +      chr1      800      901    feat
chr1     150     500    feature3      0      -      chr1      155      200    feat
chr1     150     500    feature3      0      -      chr1      800      901    feat
chr1      900     950    feature4      0      +      chr1      155      200    feat
chr1      900     950    feature4      0      +      chr1      800      901    feat
<BLANKLINE>
```

Note: This method returns a new bedtool instance

Note: For convenience, the file this bedtool object points to is passed as “-a”

Note: This method accepts either a bedtool or a file name as the first unnamed argument

Original BEDtools program help:

Program: windowBed (v2.12.0) Author: Aaron Quinlan (aaronquinlan@gmail.com) Summary: Examines a “window” around each feature in A and

reports all features in B that overlap the window. For each overlap the entire entry in A and B are reported.

Usage: windowBed [OPTIONS] -a <bed/gff/vcf> -b <bed/gff/vcf>

Options:

-abam	The A input file is in BAM format. Output will be BAM as well.
-ubam	Write uncompressed BAM output. Default is to write compressed BAM.
-bed	When using BAM input (-abam), write output as BED. The default is to write output in BAM when using -abam.
-w	Base pairs added upstream and downstream of each entry in A when searching for overlaps in B. - Creates symterical “windows” around A. - Default is 1000 bp. - (INTEGER)
-l	Base pairs added upstream (left of) of each entry in A when searching for overlaps in B. - Allows one to define assymterical “windows”. - Default is 1000 bp. - (INTEGER)
-r	Base pairs added downstream (right of) of each entry in A when searching for overlaps in B. - Allows one to define assymterical “windows”. - Default is 1000 bp. - (INTEGER)
-sw	Define -l and -r based on strand. For example if used, -l 500 for a negative-stranded feature will add 500 bp downstream. - Default = disabled.
-sm	Only report hits in B that overlap A on the same strand. - By default, overlaps are reported without respect to strand.
-u	Write the original A entry once if any overlaps found in B. - In other words, just report the fact ≥ 1 hit was found.
-c	For each entry in A, report the number of overlaps with B. - Reports 0 for A entries that have no overlap with B. - Overlaps restricted by -f.
-v	Only report those entries in A that have no overlaps with B. - Similar to “grep -v.”

sort () (wraps “sortBed”)

BedTool.**sort** (*args, **kwargs)
pybedtools help:

Note that chromosomes are sorted lexographically, so chr12 will come before chr9.

Example usage:

```
>>> a = pybedtools.BedTool(''
... chr9 300 400
... chr1 100 200
... chr1 1 50
... chr12 1 100
... chr9 500 600
... ''', from_string=True)
>>> print a.sort()
chr1      1      50
chr1     100     200
chr12     1      100
chr9      300     400
chr9      500     600
<BLANKLINE>
```

Note: This method returns a new bedtool instance

Note: For convenience, the file this bedtool object points to is passed as “-i”

Original BEDtools program help:

Program: sortBed (v2.12.0) Author: Aaron Quinlan (aaronquinlan@gmail.com) Summary: Sorts a feature file in various and useful ways.

Usage: sortBed [OPTIONS] -i <bed/gff/vcf>

Options:

-sizeA	Sort by feature size in ascending order.
-sizeD	Sort by feature size in descending order.
-chrThenSizeA	Sort by chrom (asc), then feature size (asc).
-chrThenSizeD	Sort by chrom (asc), then feature size (desc).
-chrThenScoreA	Sort by chrom (asc), then score (asc).
-chrThenScoreD	Sort by chrom (asc), then score (desc).

slop() (wraps “slopBed”)

BedTool.**slop**(*args, **kwargs)

pybedtools help:

Wraps slopBed, which adds bp to each feature. Returns a new BedTool object.

If *g* is a dictionary (for example, return values from `pybedtools.chromsizes()`) it will be converted to a temp file for use with slopBed. If it is a string, then it is assumed to be a filename.

Alternatively, use *genome* to indicate a pybedtools-created genome. Example usage:

```
>>> a = pybedtools.example_bedtool('a.bed')
```

Increase the size of features by 100 bp in either direction. Note that you need to specify either a dictionary of chromsizes or a filename containing chromsizes for the genome that your bed file corresponds to:

```
>>> c = a.slop(g=pybedtools.chromsizes('hg19'), b=100)
```

Grow features by 10 bp upstream and 500 bp downstream, using a genome file you already have constructed called 'hg19.genome'

First, create the file:

```
>>> fout = open('hg19.genome', 'w')
>>> chromdict = pybedtools.get_chromsizes_from_ucsc('hg19')
>>> for chrom, size in chromdict.items():
...     fout.write("%s\t%s\n" % (chrom, size[1]))
>>> fout.close()
```

Then use it:

```
>>> c = a.slop(g='hg19.genome', l=10, r=500, s=True)
```

Clean up afterwards:

```
>>> os.unlink('hg19.genome')
```

Note: This method returns a new bedtool instance

Note: For convenience, the file this bedtool object points to is passed as “-i”

Original BEDtools program help:

Program: slopBed (v2.12.0) Author: Aaron Quinlan (aaronquinlan@gmail.com) Summary: Add requested base pairs of “slop” to each feature.

Usage: slopBed [OPTIONS] -i <bed/gff/vcf> -g <genome> [-b <int> or (-l and -r)]

Options:

-b	Increase the BED/GFF/VCF entry by -b base pairs in each direction. - (Integer) or (Float, e.g. 0.1) if used with -pct.
-l	The number of base pairs to subtract from the start coordinate. - (Integer) or (Float, e.g. 0.1) if used with -pct.
-r	The number of base pairs to add to the end coordinate. - (Integer) or (Float, e.g. 0.1) if used with -pct.
-s	Define -l and -r based on strand. E.g. if used, -l 500 for a negative-stranded feature, it will add 500 bp downstream. Default = false.
-pct	Define -l and -r as a fraction of the feature's length. E.g. if used on a 1000bp feature, -l 0.50, will add 500 bp “upstream”. Default = false.

Notes:

1. Starts will be set to 0 if options would force it below 0.

(2) Ends will be set to the chromosome length if requested slop would force it above the max chrom length. (3) The genome file should tab delimited and structured as follows:

```
<chromName><TAB><chromSize>
```

For example, Human (hg19): chr1 249250621 chr2 243199373 ... chr18**gl000207**random 4262

Tips: One can use the UCSC Genome Browser’s MySQL database to extract chromosome sizes. For example, H. sapiens:

```
mysql -user=genome -host=genome-mysql.cse.ucsc.edu -A -e / "select chrom, size from hg19.chromInfo"> hg19.genome
```

flank() (wraps “flankBed”)

shuffle() (wraps “shuffleBed”)

BedTool.**shuffle**(**args, **kwargs*)

pybedtools help:

Shuffle coordinates.

Example usage:

```
>>> a = pybedtools.example_bedtool('a.bed')
>>> seed = 1 # so this test always returns the same results
>>> b = a.shuffle(genome='hg19', chrom=True, seed=seed)
>>> print b
chr1    59535036      59535135      feature1      0      +
chr1    99179023      99179123      feature2      0      +
chr1    186189051     186189401     feature3      0      -
chr1    219133189     219133239     feature4      0      +
<BLANKLINE>
```

Note: For convenience, the file this bedtool object points to is passed as “-i”

Original BEDtools program help:

Program: shuffleBed (v2.12.0) Author: Aaron Quinlan (aaronquinlan@gmail.com) Summary: Randomly permute the locations of a feature file among a genome.

Usage: shuffleBed [OPTIONS] -i <bed/gff/vcf> -g <genome>

Options:

-excl	A BED/GFF/VCF file of coordinates in which features in -i should not be placed (e.g. gaps.bed).
-incl	Instead of randomly placing features in a genome, the -incl options defines a BED/GFF/VCF file of coordinates in which features in -i should be randomly placed (e.g. genes.bed).
-chrom	Keep features in -i on the same chromosome. - By default, the chrom and position are randomly chosen.
-seed	Supply an integer seed for the shuffling. - By default, the seed is chosen automatically. - (INTEGER)

-f Maximum overlap (as a fraction of the **-i** feature) with an **-excl** feature that is tolerated before searching for a new, randomized locus. For example, **-f 0.10** allows up to 10% of a randomized feature to overlap with a given feature in the **-excl** file. **Cannot be used with **-incl** file.** - Default is 1E-9 (i.e., 1bp). - FLOAT (e.g. 0.50)

Notes:

1. The genome file should tab delimited and structured as follows: <chrom-Name><TAB><chromSize>

For example, Human (hg19): chr1 249250621 chr2 243199373 ... chr18**gl000207**random 4262

Tips: One can use the UCSC Genome Browser’s MySQL database to extract chromosome sizes. For example, H. sapiens:

```
mysql -user=genome -host=genome-mysql.cse.ucsc.edu -A -e / "select chrom, size from hg19.chromInfo" > hg19.genome
```

annotate() (wraps “**annotateBed**”)

BedTool.**annotate** (*args, **kwargs)
pybedtools help:

Annotate this BedTool with a list of other files. Example usage:

```
>>> a = pybedtools.example_bedtool('a.bed')
>>> b_fn = pybedtools.example_filename('b.bed')
>>> print a.annotate(files=b_fn)
chr1    1      100    feature1    0      +      0.000000
chr1    100    200    feature2    0      +      0.450000
chr1    150    500    feature3    0      -      0.128571
chr1    900    950    feature4    0      +      0.020000
<BLANKLINE>
```

Note: This method returns a new bedtool instance

Note: For convenience, the file this bedtool object points to is passed as “**-i**”

Original BEDtools program help:

Program: **annotateBed** (v2.12.0) Author: Aaron Quinlan (aaronquinlan@gmail.com) Summary: Annotates the depth & breadth of coverage of features from multiple files

on the intervals in **-i**.

Usage: **annotateBed** [OPTIONS] **-i** <bed/gff/vcf> **-files** FILE1 FILE2 .. FILEn

Options:

-names A list of names (one / file) to describe each file in **-i**. These names will be printed as a header line.

-counts Report the count of features in each file that overlap **-i**.

- Default is to report the fraction of **-i** covered by each file.

-both	Report the counts followed by the % coverage. - Default is to report the fraction of -i covered by each file.
-s	Force strandedness. That is, only include hits in A that overlap B on the same strand. - By default, hits are included without respect to strand.

`coverage()` (wraps “coverageBed”)

`BedTool.coverage(*args, **kwargs)`
pybedtools help:

```
>>> a = pybedtools.example_bedtool('a.bed')
>>> b = pybedtools.example_bedtool('b.bed')
>>> c = a.coverage(b)
>>> c.head(3)
chr1    155    200    feature5    0    -    2    45    45    1.0000000
chr1    800    901    feature6    0    +    1    1    101    0.0099010
```

Note: This method returns a new bedtool instance

Note: For convenience, the file this bedtool object points to is passed as “-a”

Original BEDtools program help:

Program: coverageBed (v2.12.0) Author: Aaron Quinlan (aaronquinlan@gmail.com) Summary: Returns the depth and breadth of coverage of features from A on the intervals in B.

Usage: coverageBed [OPTIONS] -a <bed/gff/vcf> -b <bed/gff/vcf>

Options:

-abam	The A input file is in BAM format.
-s	Force strandedness. That is, only include hits in A that overlap B on the same strand. - By default, hits are included without respect to strand.
-hist	Report a histogram of coverage for each feature in B as well as a summary histogram for all features in B. Output (tab delimited) after each feature in B: <ol style="list-style-type: none">1. depth2. # bases at depth3. size of B4. % of B at depth
-d	Report the depth at each position in each B feature. Positions reported are one based. Each position and depth follow the complete B feature.
-split	Treat “split” BAM or BED12 entries as distinct BED intervals. when computing coverage. For BAM files, this uses

the CIGAR “N” and “D” operations to infer the blocks for computing coverage. For BED12 files, this uses the Block-Count, BlockStarts, and BlockEnds fields (i.e., columns 10,11,12).

Default Output:

After each entry in B, reports:

1. The number of features in A that overlapped the B interval.
2. The number of bases in B that had non-zero coverage.
3. The length of the entry in B.
4. The fraction of bases in B that had non-zero coverage.

genome_coverage() (wraps “genomeCoverageBed”)

BedTool.genome_coverage(*args, **kwargs)

pybedtools help:

Calculates coverage at each position in the genome.

Use `bg=True` to have the resulting BedTool return valid BED-like features

Example usage:

```
>>> a = pybedtools.example_bedtool('x.bam')
>>> b = a.genome_coverage(ibam=a.fn, genome='dm3', bg=True)
>>> b.head(3)
chr2L    9329    9365    1
chr2L    10212   10248    1
chr2L    10255   10291    1
```

Note: This method returns a new bedtool instance

Note: For convenience, the file this bedtool object points to is passed as “-i”

Original BEDtools program help:

Program: genomeCoverageBed (v2.12.0) Authors: Aaron Quinlan (aaronquinlan@gmail.com)

Assaf Gordon, CSHL

Summary: Compute the coverage of a feature file among a genome.

Usage: genomeCoverageBed [OPTIONS] -i <bed/gff/vcf> -g <genome>

Options:

-ibam	The input file is in BAM format. Note: BAM must be sorted by position
-d	Report the depth at each genome position. Default behavior is to report a histogram.
-bg	Report depth in BedGraph format. For details, see: genome.ucsc.edu/goldenPath/help/bedgraph.html

-bga Report depth in BedGraph format, as above (-bg). However with this option, regions with zero coverage are also reported. This allows one to quickly extract all regions of a genome with 0 coverage by applying: “grep -w 0\$” to the output.

-split Treat “split” BAM or BED12 entries as distinct BED intervals. when computing coverage. For BAM files, this uses the CIGAR “N” and “D” operations to infer the blocks for computing coverage. For BED12 files, this uses the Block-Count, BlockStarts, and BlockEnds fields (i.e., columns 10,11,12).

-strand Calculate coverage of intervals from a specific strand. With BED files, requires at least 6 columns (strand is column 6). - (STRING): can be + or -

-max Combine all positions with a depth \geq max into a single bin in the histogram. Irrelevant for -d and -bedGraph - (INTEGER)

Notes:

1. The genome file should tab delimited and structured as follows: <chrom-Name><TAB><chromSize>

For example, Human (hg19): chr1 249250621 chr2 243199373 ... chr18**gl000207**random 4262

2. The input BED (-i) file must be grouped by chromosome. A simple “sort -k 1,1 <BED> > <BED>.sorted” will suffice.
3. The input BAM (-ibam) file must be sorted by position. A “samtools sort <BAM>” should suffice.

Tips: One can use the UCSC Genome Browser’s MySQL database to extract chromosome sizes. For example, H. sapiens:

```
mysql -user=genome -host=genome-mysql.cse.ucsc.edu -A -e / "select chrom, size from hg19.chromInfo" > hg19.genome
```

Methods for converting between formats

bed6 () (wraps “Bed12To6”)

BedTool.**bed6** (*args, **kwargs)
pybedtools help:

convert a BED12 to a BED6 file

Note: This method returns a new bedtool instance

Note: For convenience, the file this bedtool object points to is passed as “-i”

Note: This method accepts either a bedtool or a file name as the first unnamed argument

Original BEDtools program help:

Program: bed12ToBed6 (v2.12.0) Author: Aaron Quinlan (aaronquinlan@gmail.com) Summary: Splits BED12 features into discrete BED6 features.

Usage: bed12ToBed6 [OPTIONS] -i <bed12>

Options:

-n	Force the score to be the (1-based) block number from the BED12.
-----------	--

Methods for working with sequences

`sequence()` (wraps “`fastaFromBed`”)

`BedTool.sequence(fi, **kwargs)`

pybedtools help:

Wraps `fastaFromBed`. *fi* is passed in by the user; *bed* is automatically passed in as the bedfile of this object; *fo* by default is a temp file. Use `save_seqs()` to save as a file.

The end result is that this `BedTool` will have an attribute, `self.seqfn`, that points to the new fasta file.

Example usage:

```
>>> a = pybedtools.BedTool("""
... chr1 1 10
... chr1 50 55""", from_string=True)
>>> fasta = pybedtools.example_filename('test.fa')
>>> a = a.sequence(fi=fasta)
>>> print open(a.seqfn).read()
>chr1:1-10
GATGAGTCT
>chr1:50-55
CCATC
<BLANKLINE>
```

Note: This method returns a new bedtool instance

Note: For convenience, the file this bedtool object points to is passed as “-bed”

Original BEDtools program help:

Program: fastaFromBed (v2.12.0) Author: Aaron Quinlan (aaronquinlan@gmail.com) Summary: Extract DNA sequences into a fasta file based on feature coordinates.

Usage: fastaFromBed [OPTIONS] -fi <fasta> -bed <bed/gff/vcf> -fo <fasta>

Options:

-fi	Input FASTA file
-bed	BED/GFF/VCF file of ranges to extract from -fi
-fo	Output file (can be FASTA or TAB-delimited)

-name	Use the name field for the FASTA header
-tab	Write output in TAB delimited format. - Default is FASTA format.
-s	Force strandedness. If the feature occupies the antisense strand, the sequence will be reverse complemented. - By default, strand information is ignored.

mask_fasta() (wraps “maskFastaFromBed”)

BedTool.**mask_fasta**(**args, **kwargs*)

pybedtools help:

Masks a fasta file at the positions in a BED file and saves result as *out*. This method returns None, and sets self.seqfn to *out*.

```
>>> a = pybedtools.BedTool('chr1 100 110', from_string=True)
>>> fasta_fn = pybedtools.example_filename('test.fa')
>>> a = a.mask_fasta(fi=fasta_fn, fo='masked.fa.example')
>>> b = a.slop(b=2, genome='hg19')
>>> b = b.sequence(a.seqfn)
>>> print b.print_sequence()
>chr1:98-112
TTNNNNNNNNNNAT
<BLANKLINE>
>>> os.unlink('masked.fa.example')
>>> if os.path.exists('masked.fa.example.fai'):
...     os.unlink('masked.fa.example.fai')
```

Note: This method returns a new bedtool instance

Note: For convenience, the file this bedtool object points to is passed as “-bed”

Original BEDtools program help:

Program: maskFastaFromBed (v2.12.0) Author: Aaron Quinlan (aaronquinlan@gmail.com) Summary: Mask a fasta file based on feature coordinates.

Usage: maskFastaFromBed [OPTIONS] -fi <fasta> -out <fasta> -bed <bed/gff/vcf>

Options:

-fi	Input FASTA file
-bed	BED/GFF/VCF file of ranges to mask in -fi
-fo	Output FASTA file
-soft	Enforce “soft” masking. That is, instead of masking with Ns, mask with lower-case bases.
-mc	Replace masking character. That is, instead of masking with Ns, use another character.

2.5.3 BedTool methods unique to pybedtools

Introspection

`count()`

`BedTool.count()`

Number of features in BED file. Does the same thing as `len(self)`, which actually just calls this method.

Only counts the actual features. Ignores any track lines, browser lines, lines starting with a “#”, or blank lines.

Example usage:

```
a = BedTool('in.bed')
a.count()
```

`print_sequence()`

`BedTool.print_sequence()`

Print the sequence that was retrieved by the `BedTool.sequence()` method.

See usage example in `BedTool.sequence()`.

`field_count()`

`BedTool.field_count(n=10)`

Return the number of fields in the features this file contains. Checks the first *n* features.

Saving

`saveas()`

`BedTool.saveas(fn, trackline=None)`

Save BED file as a new file, adding the optional *trackline* to the beginning.

Returns a new `BedTool` for the newly saved file.

A newline is automatically added to the trackline if it does not already have one.

Example usage:

```
>>> a = pybedtools.example_bedtool('a.bed')
>>> b = a.saveas('other.bed')
>>> b.fn
'other.bed'
>>> print b == a
True

>>> b = a.saveas('other.bed', trackline="name='test run' color=0,55,0")
>>> open(b.fn).readline()
"name='test run' color=0,55,0\n"
```

Note: This method returns a new `bedtool` instance

save_seqs()

`BedTool.save_seqs(fn)`

Save sequences of features in this `BedTool` object as a fasta file *fn*.

In order to use this function, you need to have called the `BedTool.sequence()` method.

A new `BedTool` object is returned which references the newly saved file.

Example usage:

```
a = BedTool('in.bed')

# specify the filename of the genome in fasta format
a.sequence('data/genomes/genome.fa')

# use this method to save the seqs that correspond to the features
# in "a"
a.save_seqs('seqs.fa')
```

Utilities

with_attrs()

`BedTool.with_attrs(**kwargs)`

Given arbitrary keyword arguments, turns the keys and values into attributes. Useful for labeling `BedTools` at creation time.

Example usage:

```
>>> # add a "label" attribute to each BedTool
>>> a = pybedtools.example_bedtool('a.bed')
>>> b = pybedtools.example_bedtool('b.bed')
>>> for i in [a, b]:
...     print i.count(), 'features for', i.label
4 features for transcription factor 1
2 features for transcription factor 2
```

Note: This method returns a new `bedtool` instance

cat()

`BedTool.cat(other, postmerge=True, **kwargs)`

Concatenates two `BedTool` objects (or an object and a file) and does an optional post-merge of the features.

Use `postmerge=False` if you want to keep features separate.

TODO:

currently truncates at BED3 format!

`kwargs` are sent to `BedTool.merge()`.

Example usage:


```
>>> a = pybedtools.example_bedtool('a.bed')
>>> b = pybedtools.example_bedtool('b.bed')
>>> print a.cat(b)
chr1    1      500
chr1    800    950
<BLANKLINE>
```

Note: This method returns a new bedtool instance

Note: This method accepts either a bedtool or a file name as the first unnamed argument

`total_coverage()`

`BedTool.total_coverage()`

Returns the total number of bases covered by this BED file. Does a `self.merge()` first to remove potentially multiple-counting bases.

Example usage:

```
>>> a = pybedtools.example_bedtool('a.bed')
```

This does a `merge()` first, so this is what the total coverage is counting:

```
>>> print a.merge()
chr1    1      500
chr1    900    950
<BLANKLINE>

>>> print a.total_coverage()
549
```

`delete_temporary_history()`

`BedTool.delete_temporary_history(ask=True, raw_input_func=None)`

Use at your own risk! This method will delete temp files. You will be prompted for deletion of files unless you specify `ask=False`.

Deletes all temporary files created during the history of this `BedTool` up to but not including the file this current `BedTool` points to.

Any filenames that are in the history and have the following pattern will be deleted:

```
<TEMP_DIR>/pybedtools.*.tmp
```

(where `<TEMP_DIR>` is the result from `get_tmpdir()` and is by default `"/tmp"`)

Any files that don't have this format will be left alone.

(`raw_input_func` is used for testing)

Feature-by-feature operations

`each()`

`BedTool.each(func, *args, **kwargs)`

Applies user-defined function *func* to each feature. *func* must accept an `Interval` as its first argument; *args* and ***kwargs* will be passed to **func*.

func must return an `Interval` object.

```
>>> def truncate_feature(feature, limit=0):
...     feature.score = str(len(feature))
...     if len(feature) > limit:
...         feature.stop = feature.start + limit
...         feature.name = feature.name + '.short'
...     return feature

>>> a = pybedtools.example_bedtool('a.bed')
>>> b = a.each(truncate_feature, limit=100)
>>> print b
chr1    1      100    feature1    99      +
chr1    100    200    feature2    100     +
chr1    150    250    feature3.short 350     -
chr1    900    950    feature4    50      +
<BLANKLINE>
```

`filter()`

`BedTool.filter(func, *args, **kwargs)`

Takes a function *func* that is called for each feature in the `BedTool` object and returns only those for which the function returns `True`.

args and ***kwargs* are passed directly to **func*.

Returns a streaming `BedTool`; if you want the filename then use the `.saveas()` method.

```
>>> a = pybedtools.example_bedtool('a.bed')
>>> subset = a.filter(lambda b: b.chrom == 'chr1' and b.start < 150)
>>> len(a), len(subset)
(4, 2)
```

so it has extracted 2 records from the original 4.

`cut()`

`BedTool.cut(indexes)`

Similar to unix `cut` except indexes are 0-based, must be a list and the columns are returned in the order requested.

In addition, indexes can contain keys of the GFF/GTF attributes, in which case the values are returned. e.g. 'gene_name' will return the corresponding name from a GTF, or 'start' will return the start attribute of a BED Interval.

See `.with_column()` if you need to do more complex operations.

features()

`BedTool.features()`

Returns an iterator of feature objects.

Randomization helpers

randomintersection()

`BedTool.randomintersection(other, iterations, intersect_kwargs=None, shuffle_kwargs=None, debug=False)`

Performs *iterations* shufflings of self, each time intersecting with *other*.

Returns a generator of integers where each integer is the number of intersections of a shuffled file with *other*. This distribution can be used in downstream analysis for things like empirical p-values.

intersect_kwargs and *shuffle_kwargs* are passed to `self.intersect()` and `self.shuffle()` respectively. By default for `intersect`, `u=True` is specified – but `s=True` might be a useful option for strand-specific work.

Useful kwargs for *shuffle_kwargs* are `chrom`, `excl`, or `incl`. If you use the “seed” kwarg, that seed will be used *each* time `shuffleBed` is called – so all your randomization results will be identical for each iteration. To get around this and to allow for tests, `debug=True` will set the seed to the iteration number.

Example usage:

```
>>> chromsizes = {'chr1':(0, 1000)}
>>> a = pybedtools.example_bedtool('a.bed').set_chromsizes(chromsizes)
>>> b = pybedtools.example_bedtool('b.bed')
>>> results = a.randomintersection(b, 10, debug=True)
>>> print list(results)
[2, 2, 2, 0, 2, 3, 2, 1, 2, 3]
```

randomstats()

`BedTool.randomstats(other, iterations, **kwargs)`

Sends args and kwargs to `BedTool.randomintersection()` and compiles results into a dictionary with useful stats. Requires `scipy` and `numpy`.

This is one possible way of assigning significance to overlaps between two files. See, for example:

Negre N, Brown CD, Shah PK, Kheradpour P, Morrison CA, et al. 2010 A Comprehensive Map of Insulator Elements for the Drosophila Genome. PLoS Genet 6(1): e1000814. doi:10.1371/journal.pgen.1000814

Example usage:

```
Make chromsizes a very small genome for this example: >>> chromsizes = {'chr1':(1,1000)} >>> a = pybedtools.example_bedtool('a.bed').set_chromsizes(chromsizes) >>> b = pybedtools.example_bedtool('b.bed') >>> results = a.randomstats(b, 100, debug=True)
```

results is a dictionary that you can inspect. The actual overlap: `>>> print results['actual']` 3

The median of all randomized overlaps: `>>> print results['median randomized']` 2.0

The percentile of the actual overlap in the distribution of randomized overlaps, which can be used to get an empirical p-value: `>>> print results['percentile']` 90.0

INDICES AND TABLES

- *genindex*
- *modindex*
- *search*

PYTHON MODULE INDEX

p

`pybedtools`, [35](#)

INDEX

A

annotate() (pybedtools.BedTool method), 45

B

bed6() (pybedtools.BedTool method), 48

C

cat() (pybedtools.BedTool method), 52
chromsizes() (in module pybedtools), 34
chromsizes_to_file() (in module pybedtools), 34
cleanup() (in module pybedtools), 35
closest() (pybedtools.BedTool method), 39
count() (pybedtools.BedTool method), 51
coverage() (pybedtools.BedTool method), 46
cut() (pybedtools.BedTool method), 54

D

data_dir() (in module pybedtools), 34
delete_temporary_history() (pybedtools.BedTool method), 53

E

each() (pybedtools.BedTool method), 54
example_bedtool() (in module pybedtools), 34
example_filename() (in module pybedtools), 34

F

features() (pybedtools.BedTool method), 55
field_count() (pybedtools.BedTool method), 51
filter() (pybedtools.BedTool method), 54
find_tagged() (in module pybedtools), 35

G

genome_coverage() (pybedtools.BedTool method), 47
get_chromsizes_from_ucsc() (in module pybedtools), 35
get_tempdir() (in module pybedtools), 35

I

intersect() (pybedtools.BedTool method), 36
IntervalIterator() (in module pybedtools), 35

L

list_example_files() (in module pybedtools), 34

M

mask_fasta() (pybedtools.BedTool method), 50
merge() (pybedtools.BedTool method), 37

P

print_sequence() (pybedtools.BedTool method), 51
pybedtools (module), 32, 34, 35

R

randomintersection() (pybedtools.BedTool method), 55
randomstats() (pybedtools.BedTool method), 55

S

save_seqs() (pybedtools.BedTool method), 52
saveas() (pybedtools.BedTool method), 51
sequence() (pybedtools.BedTool method), 49
set_bedtools_path() (in module pybedtools), 35
set_tempdir() (in module pybedtools), 35
shuffle() (pybedtools.BedTool method), 44
slop() (pybedtools.BedTool method), 42
sort() (pybedtools.BedTool method), 41
subtract() (pybedtools.BedTool method), 38

T

total_coverage() (pybedtools.BedTool method), 53

W

window() (pybedtools.BedTool method), 40
with_attrs() (pybedtools.BedTool method), 52